

DETAILED PROJECT REPORT

Global Data Centre Industry

Growth & Financial Analysis

2020 – 2026 | Consulting-Grade Industry Report

Document Type Detailed Project Report (DPR)	Report Period 2020 – 2026 (with 2030+ Outlook)
Coverage Global (US, Europe, APAC, India)	Date April 2026

Benchmarking: Equinix · BCG · Vantage Data Centres · HPE · EdgeConneX

1. EXECUTIVE SUMMARY

The global data centre industry has undergone one of the most rapid structural transformations in the history of commercial infrastructure over the period 2020 to 2026. What began as a sector primarily serving enterprise IT hosting needs has evolved into the backbone of the digital economy — processing, transmitting, and storing the world's exploding volumes of data generated by cloud computing, artificial intelligence, the Internet of Things, streaming media, financial transactions, and an ever-expanding universe of connected devices.

This Detailed Project Report (DPR) provides a consulting-grade analysis of the global data centre industry's growth trajectory between 2020 and 2026, encompassing market sizing, regional dynamics, capital investment trends, financial modelling, company benchmarking, and a forward-looking assessment of challenges and opportunities through 2030 and beyond. The report draws on validated market intelligence from leading research firms, actual financial disclosures from publicly traded operators, and real-world project cost data from infrastructure consultancies.

1.1 Market Growth Highlights

The global data centre market — measured by total infrastructure investment including IT hardware, power systems, cooling infrastructure, and real estate — grew from approximately \$180–200 billion in 2020 to an estimated \$380–430 billion in 2025, representing a compound annual growth rate (CAGR) of approximately 13–17% depending on the methodology and scope of measurement used by different research houses. Grand View Research places the 2025 market at \$383.82 billion, projecting it to reach \$425.28 billion in 2026 and \$902.19 billion by 2033. Precedence Research similarly estimates \$386.71 billion in 2025, projecting \$430.18 billion in 2026.

These figures, while varying across methodologies, consistently signal the same macro story: the global data centre industry is one of the fastest-growing infrastructure segments in the world, with a growth rate roughly twice that of global GDP and dramatically outpacing other commercial real estate asset classes. By 2030, the industry is broadly projected to reach \$600–700 billion, with some scenarios projecting a trillion-dollar market by 2033–2035 if AI-driven demand continues its current trajectory.

\$383B+ Global Market (2025E) Grand View Research	\$425B+ Global Market (2026E) +11% YoY	11.3% Market CAGR 2025–2033 GVR Forecast	\$61B+ Record Investment (2025) New Capacity
--	---	---	---

1.2 Investment Trends

Global data centre investment reached record levels in 2025, with total new capacity investment surpassing \$61 billion — the highest single-year figure in the industry's history. Hyperscalers alone drove the lion's share of this spending: Amazon Web Services committed more than \$100 billion in total capital expenditures in 2025, Microsoft approximately \$80 billion, and Meta \$65 billion. In the United States alone, annual data centre investment was projected to exceed \$425 billion in 2025, with approximately 70% of that spending originating from hyperscale cloud providers.

The pace of U.S. data centre construction starts has been nothing short of extraordinary. Through July 2025, cumulative year-to-date construction starts reached \$26.9 billion, nearly triple the comparable period in 2024. In July 2025 alone, new construction starts totaled \$14 billion — more than double the previous single-month record of \$7.4 billion set in August 2024. The 12-month trailing average cost per data centre reached \$220 million, the third time in industry history this threshold has been crossed.

Private capital has also surged into the sector. In 2024, Vantage Data Centres secured more than \$13 billion in combined debt and equity investments from institutional backers including DigitalBridge and Silver Lake Partners, bringing its total planned and operational capacity to more than 2.6 gigawatts globally. In August 2025, Vantage announced its largest-ever commitment: over \$25 billion to develop a 1.4 GW campus in Shackelford County, Texas.

1.3 Key Financial Insights

The financial profile of data centre investments is characterized by high upfront capital requirements offset by long-term contracted cash flows, strong EBITDA margins, and rising return profiles as scale advantages compound. The following financial benchmarks represent the industry's current state:

Financial Metric	Range / Value	Notes
CapEx per MW (Standard)	\$10M – \$12M	2025–2026 baseline; US/Europe
CapEx per MW (AI-Ready)	\$15M – \$20M+	High-density GPU; liquid cooling
CapEx per MW (Hyperscale)	\$7M – \$10M	Large-scale builds; procurement leverage
Annual OpEx (30MW Facility)	\$100M+	Power + labor + maintenance
Power as % of OpEx	40% – 60%	Varies by PUE and energy market
EBITDA Margin (Colocation)	45% – 55%	Mature tier-1 operators
EBITDA Margin (Equinix 2026E)	~51%	Company guidance Feb 2026
Payback Period	8 – 12 years	Standard enterprise/colo builds
Typical IRR	10% – 15%	Hyperscale pre-leased; up to 20%+ value-add

Equinix, the world's largest data centre and colocation operator by interconnection capacity, serves as the industry's most closely watched financial bellwether. For full year 2025, Equinix reported revenue of \$9.217 billion, a 5% increase over the prior year on an as-reported basis. For 2026, the company guided to revenues of \$10.123 to \$10.223 billion with an adjusted EBITDA margin of approximately 51%, representing a ~200 basis-point margin expansion over 2025. These figures confirm the operating leverage thesis: as data centres scale, fixed cost bases are absorbed over greater revenue, driving margin expansion.

The edge data centre sub-segment — smaller distributed facilities deployed close to end users — grew from approximately \$7.2 billion in 2021 to an estimated \$14.7 – \$19.5 billion in 2025–2026, representing CAGRs of 17–25% depending on the research source. Multiple forecasters project the edge market to reach \$50–110 billion by 2030. This explosive growth is driven by 5G network rollouts, the proliferation of IoT devices, autonomous systems, and the need to process AI inference workloads at the network edge with sub-millisecond latency.

2. INDUSTRY OVERVIEW (2020–2026)

2.1 Market Size Evolution

The global data centre industry's growth story over the 2020–2026 period is one of structural acceleration driven by a convergence of technology trends, macroeconomic forces, and shifting enterprise IT strategies. The COVID-19 pandemic, which began disrupting economies in early 2020, paradoxically served as an accelerant for data centre demand: mass remote working, the surge in video conferencing, the acceleration of e-commerce, and the rapid digitization of previously offline services all drove immediate and sustained increases in cloud computing and data infrastructure demand.

From a baseline of approximately \$180–200 billion in 2020, the global market grew through successive waves of investment. The first wave (2020–2021) was driven by COVID-induced digital adoption and the maturation of hyperscale cloud platforms. The second wave (2022–2023) was characterized by enterprise cloud migration at scale and the buildout of edge infrastructure alongside 5G networks. The third and most powerful wave (2024–2026) has been dominated by artificial intelligence infrastructure, particularly the demand for GPU-dense data centres capable of training and serving large language models and other AI workloads.

Year	Est. Global Market Size	Key Driver	YoY Growth
2020	~\$180–200B	COVID-19 digital surge; cloud migration	~8–10%
2021	~\$210–230B	Hyperscale expansion; 5G early deployments	~10–12%
2022	~\$240–265B	Enterprise cloud; edge computing emergence	~12%
2023	~\$270–300B	AI infrastructure early wave; colocation growth	~12–14%
2024	~\$320–390B	AI/GenAI boom; hyperscale CapEx surge	~15–20%
2025E	~\$384–430B	Record hyperscale investment; AI data centres	~13–17%
2026E	~\$425–480B	AI inference scale-out; edge + hyperscale	~10–12%

The variation in market sizing across research firms reflects methodological differences: some measure only colocation and hyperscale revenues, others include all data centre hardware, software, and

services. The consensus directional view is unambiguous — the industry roughly doubled in total economic value between 2020 and 2026, representing a CAGR of approximately 11–16%.

2.2 CAGR Analysis

The compound annual growth rates observed across the data centre industry vary significantly by sub-segment, providing important investment and strategic insights:

Segment	CAGR (2025–2030/2033)	Primary Growth Driver
Overall Data Centre Market	10.6% – 11.5%	AI, cloud, digital transformation
Edge Data Centre Market	16.5% – 28.9%	5G, IoT, AI inference, latency
Hyperscale Data Centres	15%+ (capacity)	Cloud hyperscalers; AI training
Colocation Services	9% – 12%	Enterprise hybrid cloud
Data Centre Networking	17.6%	AI fabric; high-speed interconnect
Liquid Cooling Systems	25%+	High-density AI/GPU workloads
Edge IT Infrastructure	22.1%	Distributed AI, smart cities

The differentiated CAGR profile across sub-segments reveals a critical strategic insight: the highest growth rates are concentrated in the portions of the industry most directly exposed to AI workloads and edge deployment. Standard colocation and general compute infrastructure, while still growing solidly, are being outpaced by AI-ready and edge-specific segments. This bifurcation is reshaping capital allocation decisions, with operators investing heavily to retrofit existing facilities with liquid cooling and high-density power distribution while simultaneously building purpose-designed AI data centres from the ground up.

2.3 Regional Analysis

2.3.1 North America

North America is and remains the world's dominant data centre market, accounting for approximately 38–42% of global market revenue in 2024–2025 and holding 54% of all hyperscale data centres worldwide. The United States, as the home market of Amazon Web Services, Microsoft Azure, Google Cloud, and Meta's infrastructure, drives the majority of this demand. The U.S. data centre market was estimated at \$134.77 billion in 2025 and is projected to reach \$393.14 billion by 2035, growing at a CAGR of 11.3%.

Within North America, the key hyperscale markets of Northern Virginia (the largest data centre market globally), Silicon Valley, Dallas-Fort Worth, Phoenix, Chicago, and Atlanta continue to attract the largest share of investment. However, power constraints in established markets have driven development

toward emerging markets including Columbus (Ohio), Omaha, Reno, and secondary cities across the Southeast and Midwest. The average data centre land transaction size in the U.S. grew to 224 acres in 2024, up 144% since 2022, reflecting the shift to massive hyperscale campus developments.

U.S. data centre construction starts more than tripled year-over-year through July 2025, with the July 2025 monthly figure of \$14 billion setting an all-time single-month record. Power backup infrastructure alone accounted for 57.1% of U.S. data centre construction spending in 2024, reflecting the central challenge of securing adequate electrical power in constrained grid markets. Transformer lead times — averaging 40 weeks before 2020 — have stretched to up to four years for large models as of 2025, creating significant project delivery risks.

2.3.2 Europe

Europe represents approximately 15% of global hyperscale data centre capacity and is the world's second-largest data centre market by investment. The FLAP-D markets — Frankfurt, London, Amsterdam, Paris, and Dublin — have historically dominated European data centre activity. However, these markets face acute constraints including limited land availability, power grid saturation, strict sustainability regulations, and complex permitting processes.

The European Union's regulatory environment, including the GDPR data residency requirements, is simultaneously a growth driver (forcing local deployments) and a compliance cost for operators. The EU Energy Efficiency Directive and National Data Centre Strategies across Germany, France, and Spain are pushing operators toward renewable energy integration and green data centre certifications. European markets led globally in renewable energy as a share of total power production, making them attractive for ESG-conscious hyperscale tenants seeking to meet carbon neutrality commitments.

Emerging European markets including Madrid, Milan, Warsaw, Stockholm, and Helsinki are growing rapidly as operators seek alternatives to the saturated FLAP-D markets. The Nordic countries, with abundant hydroelectric power and natural cooling advantages, are particularly attractive for energy-intensive AI data centres. EdgeConneX, one of the benchmarked companies in this report, has been a significant driver of edge data centre development across secondary European markets.

2.3.3 Asia-Pacific

Asia-Pacific is the fastest-growing major data centre market globally, with its total operational capacity reaching 12.2 GW by the end of 2024 — a 1.6 GW increase in a single year. The regional development pipeline remains extremely strong, with an additional 14.4 GW of capacity under construction or in planning phases as of early 2025. Asia Pacific is projected to grow at a 20% CAGR through 2035 for edge data centres, the highest rate globally.

China, despite geopolitical headwinds for Western operators, accounts for 16% of global hyperscale data centres, the second-largest concentration after the U.S. Japanese markets, including Tokyo and Osaka, have seen significant investment from both domestic operators and international hyperscalers.

Singapore, historically a key APAC hub, has faced a data centre moratorium since 2019 due to power constraints, redirecting investment to alternative markets including Malaysia (Johor Bahru and Kuala Lumpur), Indonesia (Jakarta), and Thailand.

Australia is emerging as a high-growth market, driven by sovereign cloud requirements, growing digital economy demands, and increasing hyperscaler presence from AWS, Azure, and Google Cloud. South Korea has also become a significant market, with domestic technology giants Naver and Kakao building hyperscale infrastructure to serve the Korean digital economy.

2.3.4 India

India represents one of the most compelling data centre growth opportunities globally. The Indian data centre market was estimated at \$11.49 billion in 2026 and is projected to grow at the highest CAGR among all major national markets through 2030. Capacity in India is expected to quadruple by 2026 from 2022 baseline levels. Key growth markets include Mumbai (India's primary financial and technology hub), Chennai, Hyderabad, Pune, and Delhi-NCR.

Government initiatives including the National Informatics Centre's network of National Data Centres (NDCs), the Digital India program, and state-level data centre policies in Maharashtra, Telangana, and Tamil Nadu are driving a regulatory environment favorable to investment. Data localization requirements under emerging Personal Data Protection frameworks are creating additional demand for domestic data centre capacity.

India's growth in Tier 2 cities — Kochi, Ahmedabad, Kolkata, Nagpur — is particularly notable as telecom operators roll out 5G networks and as hyperscalers seek lower-cost alternatives to land-constrained Tier 1 markets. The confluence of one of the world's largest software engineering talent pools, competitive land and labor costs, improving power infrastructure, and massive latent digital demand positions India as a data centre investment destination of global significance through 2030 and beyond.

3. GROWTH DRIVERS

3.1 Artificial Intelligence & Generative AI Infrastructure Demand

No single force has reshaped the data centre industry's trajectory more profoundly than artificial intelligence — and specifically the emergence of generative AI (GenAI) following the public release of ChatGPT in November 2022. The scale of compute required to train and serve large language models (LLMs) such as GPT-4, Claude, Gemini, and Llama 3 dwarfs anything previously built in commercial data centres. A single training run for a frontier AI model can consume millions of GPU-hours and require hundreds of megawatts of sustained power delivered with exceptional reliability.

This AI infrastructure imperative has driven hyperscaler capital expenditures to unprecedented levels. Amazon Web Services, Microsoft Azure (which has invested heavily in OpenAI infrastructure), Google DeepMind's cloud infrastructure, and Meta AI Research all significantly accelerated their data centre build programs from 2023 onward. Hyperscale CapEx jumped 58% year-over-year in 2024 alone, with the trajectory continuing upward in 2025. AI-specific data centre capacity demand is forecast to be responsible for more than \$5 trillion of global data centre investment through 2030.

The technical requirements of AI workloads have also fundamentally changed data centre design parameters. Traditional enterprise data centres operated at power densities of 4–10 kilowatts per rack. Modern GPU clusters for AI training operate at 50–100+ kilowatts per rack, requiring entirely different cooling architectures. Direct liquid cooling (DLC) — where cooling liquid is piped directly to processor heat sinks — and immersion cooling (where servers are submerged in dielectric fluid) are transitioning from niche technologies to mainstream requirements for AI data centres. These cooling technologies add significant CapEx but dramatically improve energy efficiency and enable the extremely high power densities required by NVIDIA H100/H200 and B100/B200 GPU clusters.

AI Data Centre Economics

AI-ready data centres command a significant cost premium over standard facilities: \$15M–\$20M+ per MW versus \$10–\$12M for standard builds. A purpose-built 100MW AI training campus can require \$1.5–\$2B+ in total project cost. However, the revenue potential is commensurately higher: AI GPU compute commands \$2–\$10+ per GPU-hour in the spot market, while dedicated AI infrastructure leases generate revenue premiums of 30–50% over standard colocation rates.

3.2 Cloud Computing Expansion

Cloud computing has been the foundational demand driver for hyperscale data centre construction throughout the 2020–2026 period, and it remains the largest single source of data centre demand by volume. The global public cloud market has grown from approximately \$250 billion in 2020 to over \$700 billion in 2025, with Infrastructure as a Service (IaaS) and Platform as a Service (PaaS) — the components most directly tied to data centre capacity — representing the fastest-growing segments.

The three hyperscale cloud platforms — Amazon Web Services (AWS), Microsoft Azure, and Google Cloud Platform — collectively control approximately 65% of the global cloud infrastructure market. Their combined annual infrastructure capital expenditures reached approximately \$200+ billion in 2024, with significant multi-year expansion commitments in place. Each of these companies has announced global data centre expansion plans spanning dozens of new markets and hundreds of individual facilities.

Enterprise cloud migration — the systematic movement of corporate IT workloads from private on-premise infrastructure to cloud platforms — continues at pace, with most large enterprises operating in hybrid multi-cloud environments that span both public cloud and co-located private infrastructure. This hybrid model sustains demand for both hyperscale capacity (cloud) and colocation (private cloud and interconnection), creating a mutually reinforcing demand dynamic across the data centre ecosystem.

Software as a Service (SaaS) applications — including enterprise resource planning, customer relationship management, collaboration platforms, cybersecurity tools, and data analytics — have become standard corporate infrastructure, each application's growth translating directly into demand for the data centre capacity underlying its delivery. The sustained shift of enterprise software from license-based on-premise installations to subscription-based cloud delivery is a structural, multi-decade trend that continues to drive data centre expansion.

3.3 Edge Computing & 5G Network Rollouts

The global rollout of 5G wireless networks — capable of delivering data speeds up to 100 times faster than 4G and latencies below 1 millisecond — is creating an entirely new category of data centre demand. 5G networks are being designed to support applications requiring near-instantaneous data processing: autonomous vehicles that must make split-second navigation decisions, industrial robots in smart factories that respond to real-time sensor inputs, remote surgical procedures, augmented reality applications, and smart city infrastructure.

These applications cannot rely on centralized cloud data centres located hundreds or thousands of miles from end users. Even the speed of light imposes latency constraints that make round-trip communication to distant servers unacceptable for time-critical applications. Edge data centres — smaller facilities located within milliseconds of end users, typically at or near cellular base stations, in dense urban environments, or at major internet exchange points — are the infrastructure solution to this latency challenge.

The edge data centre market's explosive growth trajectory reflects this opportunity. From approximately \$7.2 billion in 2021, the edge market reached an estimated \$14.7–\$19.5 billion in 2025–2026, growing at CAGRs of 17–25%. By 2030, multiple forecasters project the edge market to reach \$50–\$109 billion, with the most aggressive projections approaching \$200 billion by 2035. EdgeConneX, one of the benchmarked operators in this report, pioneered the edge-focused deployment model and has been instrumental in developing the commercial frameworks for edge colocation.

Telecom operators globally are the primary drivers of network edge data centre deployment, leveraging their existing physical infrastructure presence at cell towers and central offices to integrate computing resources. Partnerships between telecom operators and hyperscale cloud providers — such as AT&T with Microsoft Azure, Verizon with AWS, and Deutsche Telekom with Microsoft — are creating distributed edge computing networks that deliver cloud services with far lower latency than centralized cloud alone.

3.4 Digital Transformation Across Industries

The term 'digital transformation' has become ubiquitous in corporate strategy, but its implications for data centre demand are concrete and measurable. Across virtually every industry vertical, organizations are digitizing previously analog processes, generating unprecedented volumes of data that must be stored, processed, and analyzed. Financial services firms are deploying real-time fraud detection and algorithmic trading systems. Healthcare providers are adopting AI-powered diagnostics and electronic health record systems. Retailers are building sophisticated omnichannel commerce and personalization engines. Manufacturing companies are implementing Industrial IoT sensor networks and digital twin simulations.

Each of these digital transformation initiatives creates demand for additional data centre capacity — either directly through corporate data centre investment, or indirectly through increased consumption of cloud services, SaaS applications, and managed IT services, all of which ultimately require physical data centre infrastructure. The Banking, Financial Services, and Insurance (BFSI) sector was the largest end-user segment in the global data centre market in 2024–2025, followed by IT & Telecom, Government, and Healthcare.

Government and public sector digital transformation — particularly the migration of citizen services, tax systems, healthcare records, and national security infrastructure to cloud and hybrid environments — represents a growing source of data centre demand globally. National data centre strategies in the European Union, India, Singapore, Australia, and Saudi Arabia reflect recognition of data centre infrastructure as critical national infrastructure, attracting both government investment and favorable regulatory treatment for private sector operators.

Industry Vertical	Digital Initiative	Data Centre Impact
Financial Services	Real-time trading, AI fraud detection	Ultra-low latency colo; GPU compute
Healthcare	AI diagnostics, EHR, genomics	Secure colocation; hybrid cloud
Manufacturing	Industry 4.0, IoT sensors, digital twins	Edge data centres; OT/IT integration
Retail / E-Commerce	Personalization AI, inventory management	CDN edge; cloud-native applications

Industry Vertical	Digital Initiative	Data Centre Impact
Government	Digital services, national clouds	Sovereign data centres; secure facilities
Media & Entertainment	4K/8K streaming, gaming, VR/AR	Content delivery; edge caching
Telecom	5G core, network functions virtualization	Edge data centres; carrier-neutral colo
Energy	Smart grid, predictive maintenance	Industrial edge; AI inference

4. FINANCIAL ANALYSIS

A rigorous financial analysis of data centre investments requires understanding three distinct but interrelated cost structures: Capital Expenditures (CapEx), which represent the upfront investment in building the facility; Operating Expenditures (OpEx), which represent the recurring costs of running the facility; and the Revenue Model that determines how and at what price capacity is monetized. Together, these determine the profitability profile and investment return characteristics of data centre assets.

4A. Construction Costs — Capital Expenditure (CapEx)

4A.1 Cost per Megawatt Overview

The standard unit of data centre cost analysis is the cost per megawatt (MW) of IT load capacity. This metric normalizes for facility size and enables comparison across different data centre types, markets, and operators. In 2025–2026, the cost per MW for standard (non-AI) data centre construction ranges from \$10–12 million in North American and European markets, with significant variation based on location, tier classification, and design specifications.

2025–2026 CapEx per MW Benchmarks

Standard Colocation / Enterprise: \$10M – \$12M per MW | Hyperscale (Experienced Builders): \$7M – \$10M per MW | AI-Ready High-Density: \$15M – \$20M+ per MW | Asian Markets (Best Practice): \$3.6M – \$6M per MW | Edge / Micro Data Centre: \$5M – \$8M per MW (per site)

These per-MW figures represent construction costs only — encompassing site civil works, structural steel, mechanical infrastructure (cooling systems), electrical infrastructure (switchgear, UPS, generators, power distribution), and the building envelope. They do not include IT hardware (servers, storage, networking), which adds substantially to Total Cost of Ownership, nor do they include land acquisition costs, which in major markets have surged significantly.

Land costs averaged \$5.59 per square foot (approximately \$244,000 per acre) globally in 2024. However, large parcels — those of 50+ acres suitable for hyperscale campus development — experienced a 23% year-over-year price increase in 2024, and the average data centre land transaction size grew to 224 acres. In Northern Virginia and Silicon Valley, land costs for prime locations can exceed \$1 million per acre, adding materially to total project costs.

4A.2 CapEx Component Breakdown

For a representative 100 MW hyperscale data centre, the total project cost typically ranges from \$900 million to \$1.5 billion (at \$9–15M per MW). A standard 30 MW colocation facility would cost approximately \$300–360 million at current rates. The CapEx components break down as follows:

CapEx Component	% of Total CapEx	Cost (100MW Hyperscale)	Description
IT Hardware (Servers, Storage, Networking)	35% – 45%	\$315M – \$675M	Compute nodes, SAN, spine-leaf fabric
Electrical Infrastructure	20% – 25%	\$180M – \$375M	UPS, generators, switchgear, PDUs
Cooling Systems (HVAC / Liquid)	15% – 20%	\$135M – \$300M	CRAC/CRAH, chillers, cooling towers, DLC
Core & Shell (Building)	10% – 15%	\$90M – \$225M	Structural, cladding, raised floor
Land Acquisition	5% – 10%	\$45M – \$150M	Varies by market; campus land
Power Infrastructure & Grid Connection	8% – 12%	\$72M – \$180M	Substations, utility interconnection
Security, Fire & Safety Systems	2% – 4%	\$18M – \$60M	Biometric access, CCTV, FM200
Engineering, Design & Commissioning	3% – 5%	\$27M – \$75M	Project management, validation
TOTAL	100%	\$900M – \$1.5B+	100MW hyperscale facility

For AI-specific data centres equipped with NVIDIA H100/H200 or Blackwell GPU clusters, the CapEx profile shifts dramatically. The cooling infrastructure cost increases to 25–30% of total project cost (versus 15–20% for standard builds) due to the requirement for direct liquid cooling loops capable of dissipating 50–100 kW per rack. Power infrastructure costs also rise proportionally to support the higher power densities. AI-ready facilities can cost \$1.5–2 billion for a 100 MW campus — and some of the most advanced GPU clusters being planned in 2025–2026 exceed 1 GW in total planned capacity.

4A.3 Hyperscale vs. Edge: A Cost Comparison

Parameter	Hyperscale Campus	Edge / Metro Data Centre
Typical Capacity	50 MW – 1+ GW	0.5 MW – 5 MW
CapEx per MW	\$7M – \$12M	\$5M – \$8M
Total Project Cost	\$350M – \$5B+	\$2.5M – \$40M
Land Requirement	100 – 1,000+ acres	0.5 – 5 acres
Build Timeline	18 – 36 months	6 – 18 months
Power Density per Rack	10 – 100 kW	5 – 30 kW
Cooling Type	Air + Liquid DLC / Immersion	Air / Hybrid
Tier Classification	Tier III / IV	Tier II / III

Parameter	Hyperscale Campus	Edge / Metro Data Centre
Primary Customer	Hyperscaler / Cloud (wholesale)	Enterprise / Telco / CDN
Revenue Model	Long-term wholesale lease	Colocation / MRC / Edge-as-a-service

4B. Operating Costs — Operating Expenditure (OpEx)

4B.1 OpEx Structure Overview

For a representative 30 MW data centre generating approximately \$100 million in annual revenues at a 10% target IRR, annual operating expenditures are in the range of \$60–75 million, producing an EBITDA margin of 25–40% for a relatively new facility, expanding to 45–55% as the facility matures and fixed costs are absorbed over growing revenue. The OpEx structure is dominated by power consumption, which represents the largest single cost component at 40–60% of total operating costs.

OpEx Category	% of Total OpEx	Annual Cost (30MW Facility)	Notes
Power / Electricity	40% – 60%	\$24M – \$45M	PUE-dependent; electricity rate varies
Cooling Systems (incremental)	10% – 15%	\$6M – \$11M	Chillers, cooling tower chemicals, water
Staffing & Labor	10% – 15%	\$6M – \$11M	NOC, security, facilities management
Maintenance (Hardware & Systems)	15% – 20%	\$9M – \$15M	Preventive + reactive; SLA-driven
Network & Connectivity	5% – 10%	\$3M – \$7.5M	Cross-connects, dark fiber, peering
Insurance & Property Tax	3% – 5%	\$1.8M – \$3.75M	Asset value, liability, business interruption
G&A and Corporate Overhead	3% – 5%	\$1.8M – \$3.75M	Finance, legal, compliance
TOTAL OpEx	100%	\$53M – \$97M	For 30MW colocation facility

4B.2 Power Usage Effectiveness (PUE) Analysis

Power Usage Effectiveness (PUE) is the foundational efficiency metric of data centre operations. Defined as the ratio of total facility energy consumption to the energy consumed solely by IT equipment, a PUE of 1.0 would represent perfect efficiency (all power goes to IT). In practice, the energy required for cooling, lighting, power conversion losses, and other overhead means real-world PUE always exceeds 1.0.

PUE Level	Classification	Typical Operator	Annual Overhead Power Cost (30MW)
1.0 – 1.1	World-Class	Google custom; Facebook	\$2.4M – \$4.8M overhead
1.1 – 1.2	Excellent	Major hyperscale; new builds	\$4.8M – \$9.6M overhead
1.2 – 1.5	Good / Standard	Tier-1 colo (Equinix, Digital Realty)	\$9.6M – \$24M overhead
1.5 – 2.0	Average	Enterprise; older builds	\$24M – \$48M overhead
2.0+	Inefficient	Legacy; poor cooling design	\$48M+ overhead

The energy cost modelling for a 30 MW data centre illustrates the financial materiality of PUE improvements. Assuming an electricity rate of \$0.07/kWh (industrial rate in favorable U.S. markets) and 8,760 hours of annual operation: a PUE of 1.5 means total power draw is 45 MW, costing approximately \$27.6 million annually. Improving PUE to 1.2 reduces total power draw to 36 MW, costing \$22.1 million annually — a \$5.5 million annual saving. At a data centre with a 15-year operating life, a 0.3 PUE improvement saves over \$80 million in cumulative power costs, making energy efficiency investments highly economically justified.

Modern data centres, particularly those built for AI workloads with direct liquid cooling, are achieving PUE values below 1.2. Hyperscale operators like Google and Meta report average global PUEs in the 1.06–1.12 range for their newest facilities. The segment of data centres with PUE below 1.2 is expected to grow at the fastest CAGR through 2033 as operators optimize for energy cost and ESG metrics.

4B.3 Energy Cost Modelling

Energy cost modelling for a data centre investment must account for several variables: the electricity rate negotiated with utilities or sourced from power purchase agreements (PPAs), the PUE of the facility, the IT load utilization rate, and the mix of renewable versus grid power. The following model illustrates annual electricity costs for different facility configurations:

Scenario	IT Load (MW)	PUE	Total Power (MW)	Rate (\$/kWh)	Annual Cost
Enterprise (Legacy)	10	1.8	18	\$0.09	\$14.2M
Colocation (Standard)	30	1.5	45	\$0.07	\$27.6M
Hyperscale (Efficient)	100	1.2	120	\$0.05	\$52.6M

Scenario	IT Load (MW)	PUE	Total Power (MW)	Rate (\$/kWh)	Annual Cost
AI Data Centre (DLC)	50	1.15	57.5	\$0.06	\$30.2M
Edge (Micro DC)	2	1.4	2.8	\$0.10	\$2.45M

4C. Revenue Model

4C.1 Colocation Revenue

Colocation (colo) remains the dominant revenue model for third-party data centre operators. In a colocation arrangement, the data centre operator provides the physical space, power, cooling, security, and network connectivity, while the customer supplies and manages their own IT hardware. Revenue is typically structured as a monthly recurring charge (MRC) based on the number of racks, cabinets, or square footage leased, plus power costs (sometimes passed through, sometimes bundled), and additional charges for cross-connects and value-added services.

Standard colocation rack rates in major U.S. markets range from \$800 to \$2,500+ per rack per month (RPM) for retail colocation, with higher rates in constrained markets like Silicon Valley, New York, and Singapore. Average colocation pricing has increased substantially in 2023–2026 as demand outstrips supply, with reported YoY price increases of 15–30% in the most constrained U.S. markets.

Colocation Revenue Formula

Revenue = Rack Price (\$/month) × Number of Racks × Utilization Rate × 12 months Example: 5,000 racks × \$1,200/rack/month × 85% utilization × 12 = \$61.2M/year With interconnection: +\$422/connection/month × 500,000 connections (industry scale) = Equinix interconnection revenue model

4C.2 Hyperscale Leasing (Wholesale)

Hyperscale leasing, also called wholesale data centre leasing, involves providing large blocks of data centre capacity — typically 1 MW to 100+ MW — to a single customer (cloud provider or large enterprise) under long-term contracts of 10–20 years. Revenue is structured on a per-MW basis, with typical wholesale rates ranging from \$80–150 per kW per month in mature U.S. markets (translating to \$80,000–150,000 per MW per month, or \$960,000–1,800,000 per MW per year).

The hyperscale leasing model provides highly predictable, contracted cash flows. Operators like Vantage Data Centres, CyrusOne, and QTS Realty Trust have built business models centred on pre-leasing hyperscale capacity to creditworthy cloud tenants before construction begins. This build-to-suit, pre-leased model substantially reduces development risk and enables aggressive financing structures.

Equinix's xScale joint venture program provides dedicated hyperscale capacity to the company's largest tenants while preserving interconnection monetization at the retail level.

4C.3 Interconnection Services

Interconnection services — the physical and logical connections between different networks, cloud providers, and customers within a data centre — represent one of the most valuable and highest-margin revenue streams in the data centre industry. Unlike commodity compute and storage capacity, interconnection is highly sticky (customers rarely disconnect established peering relationships), scales well with limited incremental cost, and commands premium pricing relative to underlying physical infrastructure.

Equinix, the world leader in colocation interconnection, surpassed 500,000 total interconnections globally in 2025 — the most in the industry. Interconnection revenues at Equinix reached \$422 million in Q3 2025 alone. The company's Equinix Fabric virtual interconnection platform extends this model into software-defined networking, enabling customers to establish cloud, network, and enterprise connections programmatically. For every interconnection, Equinix earns a monthly recurring fee (\$200–\$1,000+ depending on type and bandwidth), and the near-zero marginal cost of adding a connection drives extraordinary incremental margins.

4C.4 Edge Deployment Monetization

Edge data centre monetization models are still evolving, with several distinct approaches emerging in the market. The network edge colocation model replicates traditional colocation at a much smaller scale: telecom operators deploy small data centres at or near cellular base stations, leasing rack space to enterprise customers and CDN providers who need locally hosted compute. Edge-as-a-Service (EaaS) is an emerging model in which operators offer managed compute and storage at the edge on a subscription or consumption basis. Multi-access Edge Computing (MEC) platforms, standardized by ETSI, enable standardized edge application hosting on telco infrastructure.

Pricing for edge deployments typically commands a premium over centralized colocation due to the scarcity and distributed nature of edge locations. Edge rack rates can be 30–50% higher per kW than equivalent centralized colocation. The small size of edge facilities also means lower utilization rates in early deployment phases, compressing margins. As edge markets mature and fill rates improve, the unit economics are expected to become highly favorable given the premium pricing.

4D. Profitability Analysis

4D.1 EBITDA Margins

EBITDA (Earnings Before Interest, Taxes, Depreciation, and Amortization) margin is the standard profitability metric for data centre operators, reflecting the cash generation capacity of operations before the significant non-cash charges associated with large physical asset bases. Mature, stabilized data centre facilities and portfolios consistently generate EBITDA margins of 45–55%, making the sector one of the most profitable segments of commercial real estate.

Operator Type	Typical EBITDA Margin	Key Margin Driver	Example
Global Interconnection Leader	49% – 55%	Interconnection premium; network effects	Equinix (51% 2026E)
Hyperscale Wholesale	40% – 50%	Scale; long-term contracts; low churn	Digital Realty, Vantage
Edge / Regional Colo	30% – 45%	Still building; smaller scale	EdgeConneX, Cologix
Hybrid Cloud / Managed	20% – 35%	Higher service delivery costs	HPE GreenLake, IBM Cloud
Build-to-Suit Hyperscale	50% – 60%	Pre-leased; cost certainty; long tenure	Vantage xScale campuses

Equinix's financial trajectory illustrates the EBITDA margin expansion story beautifully. The company reported an adjusted EBITDA margin of 47% in 2024, expanding to 49% in 2025, and guided to approximately 51% for 2026 — a 200 basis-point improvement driven by operating leverage as revenue scales faster than the fixed cost base. This trajectory demonstrates the fundamental economics of mature data centre portfolios: incremental revenue largely flows through to EBITDA at very high margins once fixed infrastructure costs are covered.

4D.2 IRR, ROI, and Payback Period

Key Financial Formulas

TOTAL COST OF OWNERSHIP (TCO): $TCO = CapEx + (OpEx \times Operating\ Years) + Maintenance + Downtime\ Costs$
Example (30MW, 15 years): $\$360M + (\$70M \times 15) + \$50M + \$15M = \$1,475M$
REVENUE MODEL: $Revenue = Rack\ Price \times Racks \times Utilization\ Rate$
Example: $\$1,200 \times 5,000\ racks \times 85\% \times 12\ months = \$61.2M/year$
EBITDA: $EBITDA = Revenue - OpEx$
Example: $\$100M - \$60M = \$40M$ (40% margin, building phase)
Mature: $\$120M - \$55M = \$65M$ (54% margin, stabilized)
IRR: IRR is the discount rate at which NPV of all cash flows = 0
Typical data centre IRR: 10–15% (standard colo); 12–18% (hyperscale pre-leased); 15–20%+ (value-add)

Investment Metric	Standard Colo	Hyperscale (Pre-Leased)	AI Data Centre	Edge (Single Site)
Total CapEx	\$300M – \$400M	\$700M – \$1.5B	\$1.0B – \$2.0B+	\$5M – \$40M
Annual Revenue (Stabilized)	\$80M – \$120M	\$100M – \$200M	\$150M – \$400M	\$2M – \$15M
EBITDA Margin (Stabilized)	45% – 52%	50% – 58%	40% – 55%	30% – 45%
Payback Period	8 – 12 years	7 – 10 years	6 – 10 years	10 – 15 years

Investment Metric	Standard Colo	Hyperscale (Pre-Leased)	AI Data Centre	Edge (Single Site)
Target IRR	10% – 14%	12% – 16%	12% – 18%	8% – 14%
Lease Term (Typical)	3 – 10 years (retail)	10 – 20 years (wholesale)	5 – 15 years	1 – 5 years

The risk-return profile of data centre investments varies significantly across these categories. Pre-leased hyperscale builds offer the most favorable risk-adjusted returns: with a creditworthy anchor tenant committed under a long-term contract before construction begins, development risk is substantially mitigated, and the predictable cash flows support aggressive debt financing. AI data centres command the highest potential returns but carry greater demand risk (the AI infrastructure market is evolving rapidly) and technology risk (GPU generations turn over, potentially reducing the useful life of highly specialized infrastructure).

5. COMPANY BENCHMARKING

This section provides detailed financial and strategic analysis of five key companies that collectively represent the full spectrum of the data centre industry ecosystem: from pure-play colocation and interconnection operators to hyperscale campus developers, edge infrastructure specialists, hardware and hybrid cloud providers, and strategic consulting perspectives.

5.1 Equinix (NASDAQ: EQIX)

5.1.1 Corporate Profile

Equinix, Inc. is the world's largest data centre and interconnection company, operating as a Real Estate Investment Trust (REIT). Founded in 1998 and headquartered in Redwood City, California, Equinix operates more than 260 International Business Exchange (IBX) data centres across 72 markets in 35 countries. The company's core value proposition is the physical co-location of diverse networks, cloud providers, and enterprises in shared facilities optimized for interconnection — creating 'digital ecosystems' where participants can directly exchange traffic and data.

As of 2025, Equinix employs approximately 13,716 people globally and serves more than 10,000 customers across financial services, cloud, digital media, content, network, and enterprise verticals. The company surpassed 500,000 total interconnections globally in 2025, the most of any operator in the world, and this interconnection density creates powerful network effects: each new customer or connection adds value to all existing participants.

5.1.2 Financial Performance (2024–2026)

Metric	2024 Actual	2025 Actual	2026 Guidance
Total Revenue	\$8.748B	\$9.217B	\$10.123B – \$10.223B
Revenue Growth (YoY)	+7% (reported)	+5% (reported)	+10–11% (reported)
Adjusted EBITDA	\$4.097B	\$4.530B	\$5.141B – \$5.221B
Adjusted EBITDA Margin	47%	49%	~51%
Operating Income	~\$1.85B (2025)	\$1.848B	Growing
AFFO	Growing	Growing	\$4.158B – \$4.238B
Quarterly Dividend	Growing	\$4.69/share (Q3)	\$5.16/share (+10%)
Total CapEx (2026E)	N/A	N/A	\$3.655B – \$4.155B
Interconnection Revenues (Q3 2025)	N/A	\$422M/quarter	~\$1.7B+ annualized
Annualized Gross Bookings (2025)	N/A	\$1.6B (+27% YoY)	Accelerating

Equinix's 2025 performance demonstrated the strength of its interconnection-centric model in an AI-driven demand environment. Record annualized gross bookings of \$474 million in Q4 2025 — up 42% year-over-year — reflect surging enterprise demand for AI, cloud, and network ecosystem connectivity at Equinix's globally interconnected campuses. The 10% quarterly dividend increase to \$5.16 per share marked the company's 11th consecutive year of dividend growth, reflecting confidence in sustained cash flow generation.

5.1.3 Strategic Model & AI Positioning

Equinix's strategy is built on three interlocking pillars: physical interconnection (IBX data centres), digital interconnection (Equinix Fabric, Equinix Metal), and ecosystem development (Platform Equinix). As AI workloads have emerged as the dominant driver of enterprise IT investment, Equinix has positioned itself as the critical interconnection layer between AI inference platforms (deployed at the Equinix edge) and hyperscale AI training infrastructure.

The xScale program — Equinix's wholesale joint venture vehicle — provides hyperscale tenants like Google, Amazon, Microsoft, and Meta with large dedicated power envelopes within Equinix's global network, while Equinix retains the retail interconnection revenue from the dense ecosystem surrounding these deployments. xScale has allowed Equinix to capture the growth in hyperscale demand without dramatically shifting its balance sheet risk profile, as joint venture partners provide the majority of construction financing.

Land acquisitions across Amsterdam, Chicago, Johannesburg, London, and Toronto metros announced in Q3 2025 — capable of supporting over 900 MW of retail and xScale capacity at full build-out — signal continued long-term expansion. With an adjusted EBITDA margin trajectory pointing toward 55%+ in the next several years, Equinix is entering a period of accelerating cash flow generation as its invested capital base reaches maturity.

5.2 Vantage Data Centres

5.2.1 Corporate Profile

Vantage Data Centres is one of the world's largest private hyperscale data centre developers, operating primarily in North America and Europe. Founded in 2010 and headquartered in Denver, Colorado, Vantage focuses on building massive multi-building campus developments designed specifically for hyperscale cloud providers and AI infrastructure tenants. Unlike Equinix's retail/interconnection model, Vantage operates entirely in the wholesale hyperscale segment, building facilities of 20 MW to 1.4 GW for a single or small number of anchor tenants.

Vantage is backed by leading institutional investors including DigitalBridge (a leading digital infrastructure investment firm) and Silver Lake Partners. The company's financing model has been pioneering: Vantage developed the first data centre asset-backed securitization (ABS) in North

America, the first data centre ABS in EMEA, and in June 2025, the first euro-denominated data centre ABS in Continental Europe (raising €640 million / \$744 million in securitized term notes).

5.2.2 Financial Scale & Investment Activity

Metric	Value	Year / Context
Total Debt & Equity Raised (2024)	\$13B+	New and existing investors
Planned / Operational Global Capacity	2.6+ GW	As of end 2024
\$9.2B Equity Raise	\$9.2B	DigitalBridge & Silver Lake investment
\$3B European Portfolio Financing	\$3B	EMEA portfolio refinancing
Texas Mega-Campus Commitment	\$25B+	1.4 GW Shackelford County, TX (Aug 2025)
Euro ABS (Jun 2025)	€640M (\$744M)	First euro data centre ABS in Continental Europe

Vantage's announcement in August 2025 of a \$25+ billion commitment to develop a 1.4 GW data centre campus in Shackelford County, Texas — the company's largest-ever single investment — illustrates the extraordinary scale of AI infrastructure demand and the confidence of institutional capital in the sector's long-term growth trajectory. A 1.4 GW campus would represent more total power capacity than all but the very largest hyperscale campuses in history and would require massive grid infrastructure investments alongside the facility itself.

5.2.3 Build-to-Suit Strategy

Vantage's competitive advantage lies in its ability to execute at scale on build-to-suit hyperscale campuses. The build-to-suit model means that Vantage designs and constructs facilities to the exact specifications of anchor tenants (typically a major cloud provider), with the tenant committing to occupy the space under a long-term lease before or during construction. This pre-commitment structure eliminates speculative development risk, provides certainty of revenue, and enables the aggressive debt financing that characterizes Vantage's capital structure.

Vantage's campus strategy — grouping multiple large data centre buildings on single contiguous land parcels — delivers significant cost and operational efficiencies compared to standalone facilities. Shared substation infrastructure, fiber plant, cooling water systems, and security perimeters reduce per-MW construction and operating costs as campus scale increases. This campus model has become the dominant paradigm for hyperscale development globally, with Northern Virginia, Phoenix, and Columbus (Ohio) featuring some of the world's largest data centre campus concentrations.

5.3 EdgeConneX

5.3.1 Corporate Profile

EdgeConneX is a global data centre operator specializing in edge and regional data centre deployments, serving markets underserved by traditional large-scale data centre operators. Founded in 2013 and headquartered in Denver, Colorado, EdgeConneX pioneered the 'edge data centre' concept — smaller facilities located in secondary and tertiary markets, positioned at internet exchange points or near end-user populations to reduce latency and improve content delivery performance.

EdgeConneX operates across North America and Europe, with a distinctive strategy of entering markets before hyperscale demand fully materializes — a 'market development' approach that has allowed the company to build preferred-provider relationships with major CDN providers, streaming companies, and telecom operators. The company's \$3 billion European portfolio financing (alongside Vantage's and STACK's capital raises) reflects the institutional confidence in its scale of operations.

5.3.2 Edge Infrastructure Model

EdgeConneX's facilities range from micro data centres of sub-1 MW in small markets to 10–20 MW regional data centres in secondary cities. The company's approach recognizes that the global internet is not monolithic: content delivery, streaming, and latency-sensitive applications require distributed infrastructure at 'the edge' of the network, close to the populations consuming that content.

EdgeConneX's customers include major CDN providers (Cloudflare, Fastly, Akamai), streaming services (Netflix, Disney+), telecom operators, gaming companies, and enterprise cloud users requiring low-latency regional compute. Its revenue model combines traditional colocation (monthly recurring revenue per rack) with wholesale contracts for larger customers requiring dedicated power envelopes of 2–10 MW.

5.3.3 Strategic Positioning in the Edge Market

EdgeConneX Metric	Description
Market Focus	Secondary cities, internet exchange points, underserved markets
Facility Size Range	0.5 MW to 20+ MW per site
Revenue Model	Colocation MRC + wholesale edge contracts
Key Customers	CDN providers, streaming services, telecom operators
Geographic Footprint	North America + Europe (40+ markets)
Differentiation	First-mover in edge markets; long-term preferred agreements
Growth Driver	5G densification; streaming growth; enterprise edge AI

5.4 Hewlett Packard Enterprise (HPE)

5.4.1 Corporate Profile

Hewlett Packard Enterprise (HPE) occupies a unique position in the data centre ecosystem as the leading infrastructure hardware provider and a pioneer of hybrid cloud platforms. Unlike colocation operators, HPE does not build or operate data centre facilities for external customers; instead, it provides the servers, networking equipment, storage systems, and software management platforms that populate both enterprise-owned data centres and third-party colocation/hyperscale facilities.

Founded in 2015 as the enterprise-focused spinoff of the original Hewlett-Packard Company, HPE is headquartered in Spring, Texas and operates globally across financial services, healthcare, government, manufacturing, and technology sectors. HPE holds approximately 16% share of the global edge data centre infrastructure market (servers, storage, networking) and is among the top five players in the edge data centre market globally (alongside Huawei, Dell, IBM, and Schneider Electric), which collectively held approximately 28.4% market share in 2025.

5.4.2 HPE GreenLake: The Hybrid Cloud Platform

HPE's strategic response to the cloud computing era is GreenLake — a hybrid cloud platform that brings cloud-like consumption economics (pay-per-use, subscription-based pricing, elastic scaling) to on-premise and co-located infrastructure. GreenLake allows enterprise customers to run cloud-native workloads on HPE infrastructure deployed in their own data centres, removing the need to migrate sensitive or latency-critical workloads to public cloud while achieving operational simplicity comparable to public cloud.

For data centre operators, HPE represents a critical vendor ecosystem partner: HPE's ProLiant server line, Aruba networking (acquired 2015), Nimble and Primera storage systems, and Synergy composable infrastructure are components of the infrastructure stack across thousands of data centre deployments globally. As AI workloads have grown, HPE has expanded its offering to include AI-specific compute: the HPE Cray EX supercomputer platform powers some of the world's largest AI training clusters, and HPE announced a partnership with KDDI and NVIDIA to launch the Osaka Sakai AI Data Centre by early 2026, featuring NVIDIA Blackwell-accelerated infrastructure with HPE's hybrid cooling system.

5.4.3 HPE's Financial Footprint in Data Centres

HPE Metric	Data
Edge Data Centre Market Share	~16% (infrastructure hardware)
Edge Market Ranking	Top 5 globally (with Huawei, Dell, IBM, Schneider)
Key Products	ProLiant servers, Aruba networking, Cray supercomputers, GreenLake

HPE Metric	Data
Strategic AI Partnership	KDDI + NVIDIA Osaka Sakai AI DC (NVIDIA Blackwell)
Hybrid Cooling	Air + direct liquid cooling for AI workloads
Revenue Model	Hardware sales + GreenLake subscription + managed services

5.5 Boston Consulting Group (BCG)

5.5.1 BCG's Strategic Perspective on Data Centre Market

Boston Consulting Group, the global management consulting firm, is not a data centre operator or hardware provider — but its research and advisory work provides critical strategic intelligence on the data centre industry's direction, challenges, and opportunities. BCG has published extensively on AI infrastructure demand, data centre sustainability, power constraints, and the strategic choices facing both operators and investors in the sector.

BCG's research positions the data centre industry at an inflection point driven by four converging forces: the exponential growth of AI workloads, the physical limits of power infrastructure, the urgency of sustainability commitments, and the geopolitical fragmentation of digital infrastructure. Understanding BCG's strategic frameworks is valuable for any organization making investment decisions in or around the data centre sector.

5.5.2 BCG on AI-Driven Demand

BCG research has consistently highlighted that AI is not merely an incremental demand driver for data centres — it represents a step-change in the computational intensity of digital services. The training of a single frontier AI model can consume more electricity than hundreds of thousands of homes in a year. The scaling laws that have driven AI capability improvements (more data + more compute = better models) show no signs of abating, meaning AI-driven data centre demand growth is likely to remain exceptional for the foreseeable future.

BCG analysts have noted that data centre operators face a strategic choice between two models: the AI infrastructure provider (building massive GPU clusters and selling compute capacity) and the AI interconnection hub (Equinix's model — connecting the AI compute clusters, cloud platforms, and enterprise users through high-density interconnection). Both models have merit, but they require different CapEx profiles, customer relationships, and operational capabilities.

5.5.3 BCG on Capacity and Power Constraints

BCG has flagged power availability as the defining constraint on data centre growth in 2024–2030. In Northern Virginia, Phoenix, Dublin, Singapore, and other major markets, data centre operators report waiting times of 5+ years for new utility power connections — a fundamental bottleneck that no amount

of investment capital can immediately resolve. BCG's analysis suggests that solving the power constraint will require three parallel tracks: utility grid investment (requiring regulatory support), on-site power generation (natural gas, small modular reactors, fuel cells), and dramatically improved energy efficiency (through PUE reduction and liquid cooling adoption).

The consulting firm has also highlighted the talent constraint: data centre operators, hyperscalers, and developers are all competing for a limited pool of skilled professionals in data centre design, construction, operations, and mechanical/electrical engineering. This talent scarcity is extending project timelines and inflating labor costs, compounding the supply shortage in major markets.

6. COMPREHENSIVE FINANCIAL MODEL

This section presents a detailed, integrated financial model for a representative data centre investment, illustrating the complete lifecycle economics from development through stabilized operations. All formulas are explained with worked examples to ensure full transparency of assumptions and calculations.

6.1 Base Case: 30 MW Colocation Facility

Model Parameters

Capacity: 30 MW IT Load | Location: Tier-1 U.S. Market | Tier Classification: Tier III | PUE: 1.4 | Total Racks: 5,000 | Build Period: 24 months | CapEx/MW: \$11M | Utilization Target: 85% (Year 5) | Electricity Rate: \$0.07/kWh | Rack Rate: \$1,200/month | Analysis Period: 15 years

6.1.1 Capital Expenditure Build-Up

CapEx Component	Basis	Cost	% of Total
Land Acquisition (5 acres)	\$500K/acre	\$2.5M	0.8%
Site Preparation & Civil Works	Lump sum	\$15M	4.5%
Core & Shell Building (350,000 sq ft)	\$85/sq ft	\$29.75M	9.0%
Electrical Infrastructure (UPS, Generators, Switchgear)	Per MW	\$75M	22.7%
Cooling Systems (CRAC/CRAH, Chillers)	Per MW	\$54M	16.4%
Power Distribution (PDUs, Busways)	Per MW	\$24M	7.3%
IT Hardware (Servers, Network, Storage)	Per rack × 5,000	\$90M	27.3%
Security, Fire Suppression, Access Control	Lump sum	\$12M	3.6%
Engineering, Design, Project Management	5% of above	\$14.75M	4.5%
Contingency (8%)	8% of above	\$12.6M	3.8%
TOTAL CapEx		\$329.6M	100%

6.1.2 Revenue Model — Year 1 to Year 5

Revenue ramps as the facility fills. Year 1 assumes 40% utilization (industry standard for new builds), growing to 75% by Year 3 and 85% by Year 5 as sales cycles complete and anchor tenants expand.

Revenue Component	Year 1	Year 2	Year 3	Year 5 (Stabilized)
Rack Utilization Rate	40%	60%	75%	85%
Racks Occupied	2,000	3,000	3,750	4,250
Colocation Revenue (\$M)	\$28.8M	\$43.2M	\$54.0M	\$61.2M
Power Revenue (pass-through, \$M)	\$8.4M	\$12.6M	\$15.75M	\$17.85M
Cross-Connect / Interconnection (\$M)	\$2.0M	\$3.5M	\$5.0M	\$8.0M
Managed Services / VAS (\$M)	\$1.5M	\$2.5M	\$4.0M	\$6.0M
TOTAL Revenue (\$M)	\$40.7M	\$61.8M	\$78.75M	\$93.05M

6.1.3 Operating Expenditure — Stabilized Year

OpEx Category	Annual Cost	% of Revenue (Year 5)	Notes
Electricity (30MW × 1.4 PUE × 8,760 hrs × \$0.07)	\$25.7M	27.6%	Largest single cost item
Cooling Systems (chemicals, water, maintenance)	\$4.5M	4.8%	Variable with utilization
Staffing (50 FTE × \$85K avg)	\$4.25M	4.6%	NOC, ops, security, management
Hardware Maintenance & Warranty	\$6.0M	6.5%	3% of IT CapEx annually
Network & Connectivity (transit, cross-connects)	\$3.5M	3.8%	Upstream bandwidth, IX fees
Insurance & Property Tax	\$2.5M	2.7%	Asset value coverage
G&A and Corporate Overhead Allocation	\$2.0M	2.2%	Finance, legal, compliance

OpEx Category	Annual Cost	% of Revenue (Year 5)	Notes
TOTAL OpEx	\$48.45M	52.1%	Year 5 stabilized
EBITDA	\$44.6M	47.9%	Revenue (\$93.05M) – OpEx (\$48.45M)

6.1.4 Full TCO and Returns Analysis

TCO Formula Applied

$\text{TCO} = \text{CapEx} + (\text{OpEx} \times \text{Years}) + \text{Maintenance} + \text{Downtime Costs}$
 $\text{TCO (15-year)} = \$329.6\text{M} + (\$48.45\text{M} \times 15) + \$12\text{M} + \$5\text{M} = \$329.6\text{M} + \$726.75\text{M} + \$17\text{M} = \$1,073.35\text{M}$
 Cumulative Revenue (15-year, ramping then stable): ~\$1,050M – \$1,200M
 Net Return before financing: ~\$130M – \$250M over 15 years
 Levered IRR (with 60% LTV debt at 5.5%): ~12% – 14%
 Payback Period (from operations commencement): ~9 – 11 years

6.2 AI Data Centre Financial Model — Sensitivity

Scenario	GPU Rack Rate	Utilization	EBITDA Margin	Project IRR
Bear Case	\$4,000/rack/month	60%	35%	8% – 10%
Base Case	\$6,000/rack/month	75%	45%	12% – 14%
Bull Case	\$8,000/rack/month	90%	55%	16% – 20%
Hyperscale Pre-Lease	\$150/kW/month (wholesale)	95%	52%	14% – 18%

7. RISKS & CHALLENGES

7.1 Power Constraints: The Defining Bottleneck

The single most significant operational and strategic challenge facing the data centre industry in 2025–2030 is the availability of adequate electrical power. Electricity grids in the world's most desirable data centre markets — Northern Virginia, Phoenix, Silicon Valley, Dublin, Singapore, Amsterdam — are at or approaching saturation. Data centre operators seeking new power connections in these markets report utility queue wait times of 5–7 years in the most constrained cases, fundamentally limiting the pace at which new capacity can be brought online regardless of available capital.

By 2028, data centres could demand up to 130 GW of global power — nearly triple the approximately 47 GW consumed in 2024. In the United States alone, data centre growth is expected to drive 60% of the total increase in electrical load through 2030 — more than any other sector including electric vehicles. This power demand surge is creating grid investment requirements that utility companies, which have historically planned capital budgets on 20-year horizons, are struggling to accommodate on the accelerated timelines demanded by AI infrastructure development.

Transformer procurement represents a specific and acute bottleneck. Large power transformers — essential components of data centre electrical infrastructure — saw cost increases of 77% for power units and up to 95% for distribution units between 2019 and 2025. Lead times that averaged 40 weeks before 2020 now extend to 4 years for large models. Wood Mackenzie projected a 30% supply deficit for power transformers and a 10% shortfall for distribution units in 2025, with the situation expected to worsen before improving. A single large hyperscale campus can require dozens of large transformers, making procurement securing a critical path item.

Power Constraint Impact

- Northern Virginia: 5+ year utility queue for new power connections in some submarkets
- Phoenix: Moratoriums on new data centre power connections discussed by local utilities
- Singapore: Active data centre moratorium since 2019 (now partially lifted with restrictions)
- Dublin: Power constraints forcing new developments to outer Ireland/UK markets
- Amsterdam: Strict sustainability regulations limiting new data centre approvals
- Transformer lead times: Up to 4 years for large models (vs 40 weeks pre-2020)

7.2 High Capital Intensity

Data centre investment requires massive, upfront capital commitments with long payback periods, creating significant financial risk for operators. A 100 MW hyperscale campus requires \$900 million to \$1.5 billion in construction capital alone, before IT hardware and ongoing working capital. An AI-focused 500 MW campus could require \$5–10 billion in total capital. These capital requirements necessitate access to large-scale debt and equity markets, sophisticated financial structuring, and strong relationships with institutional investors.

The capital intensity creates several risk vectors: interest rate risk (data centre development projects are typically financed with construction loans converting to permanent debt; rising interest rates increase carrying costs and reduce project IRR), refinancing risk (large debt maturities must be refinanced in potentially challenging market conditions), and execution risk (cost overruns and construction delays are common in complex infrastructure projects, particularly given supply chain disruptions affecting electrical equipment). Overall construction costs rose 7–8% on an annualized basis since 2022, driven by equipment shortages, labor market tightness, and commodity price volatility.

7.3 ESG & Sustainability Pressures

The environmental footprint of data centres — particularly their energy consumption, carbon emissions, and water usage — has become an increasingly prominent ESG concern for both operators and their investors. Global data centres consumed approximately 415 TWh of electricity in 2024, equivalent to an average continuous power demand of approximately 47 GW. By 2030, this figure could reach 130 GW, raising profound questions about the compatibility of data centre growth with national and global carbon reduction targets.

Major hyperscale operators — Amazon Web Services, Google, Microsoft, and Meta — have all made ambitious sustainability commitments: carbon neutrality, 100% renewable energy matching, and water neutrality. Equinix has committed to 100% renewable energy for its global operations. However, meeting these commitments while simultaneously scaling at unprecedented rates requires massive procurement of renewable energy capacity and creates significant complexity in markets where renewable energy availability is limited.

Water consumption is an emerging sustainability challenge, particularly for air-cooled data centres that use evaporative cooling towers. A 100 MW data centre can consume millions of gallons of water annually. In water-stressed regions like Arizona and Nevada, this creates both operational risk (water availability) and community relations challenges. The transition to liquid cooling technologies — which consume substantially less water — is partly motivated by this concern alongside the energy efficiency benefits.

7.4 Underutilization Risk & Demand Uncertainty

A significant and often underappreciated risk for data centre investors is the potential for demand to fall short of projections, leaving expensive capacity underutilized. Data centre development cycles are long (18–36 months from commitment to opening) and capacity, once built, cannot be easily repurposed. If AI demand growth proves more volatile than anticipated — if large language model training becomes dramatically more efficient (reducing GPU demand per unit of output), if hyperscaler CapEx cycles retrench, or if economic conditions deteriorate — operators who have built speculatively could face sustained periods of low utilization.

The AI-specific demand risk is particularly relevant. A significant portion of current hyperscale CapEx is driven by a competitive race to secure GPU capacity and training infrastructure — a dynamic that could stabilize or reverse if AI scaling laws mature, if algorithmic efficiency improvements reduce compute requirements, or if a technological shift (such as the emergence of highly efficient neuromorphic computing) disrupts the GPU-centric AI training paradigm.

Risk Factor	Severity	Probability (5yr)	Mitigation Strategy
Power Constraints	High	High (already occurring)	Alternative markets; on-site generation; PPA
Interest Rate / Financing Risk	High	Medium	Long-term fixed debt; hedging
Construction Cost Overruns	Medium-High	Medium	Fixed-price contracts; modular construction
Demand Shortfall (AI)	Medium	Low-Medium	Pre-leasing; diversified customer base
Technology Obsolescence	Medium	Low-Medium	Flexible design; modular retrofit
Regulatory / Permitting Delays	High	High	Market diversification; government relations
Cybersecurity Breach	High	Medium	Multi-layered security; SOC 24/7
ESG / Regulatory Compliance	Medium	High (trend)	Renewable PPAs; liquid cooling; green certification
Geopolitical Risk	Medium	Medium	Geographic diversification; sovereign cloud

8. FUTURE OUTLOOK (2030 AND BEYOND)

8.1 AI Data Centres: The Next Generation

The AI data centre of 2030 will look fundamentally different from even the most advanced facilities being built in 2025–2026. As AI model sizes continue to grow and inference workloads (deploying trained models) scale to serve billions of users, the computational requirements will continue to rise. NVIDIA's Blackwell GPU architecture (B100/B200/GB200), already being deployed in 2025, is being succeeded by further generations with dramatically higher compute density and interconnect bandwidth. Each GPU generation compounds the value of purpose-designed liquid-cooled infrastructure.

The concept of the 'AI factory' — a term coined by NVIDIA CEO Jensen Huang — captures this evolution. Rather than traditional data centres optimized for diverse workloads, AI factories are specialized facilities optimized for a single purpose: running AI compute at maximum efficiency. These facilities are characterized by GPU cluster densities of 50–100+ kW per rack, ultra-high-bandwidth NVLink and InfiniBand interconnects between GPUs, direct liquid cooling loops integral to the server hardware, and power requirements that can reach 500 MW to 1 GW+ for a single campus.

Sector forecasts to 2030 consistently project the overall data centre market reaching \$600–700 billion globally, with AI infrastructure representing an increasingly dominant share. AI-specific data centre capacity demand is forecast to be responsible for more than \$5 trillion in global investment through 2030. The U.S. market alone is expected to nearly triple in size from its 2025 level to approximately \$393 billion by 2035.

8.2 Liquid Cooling: The Cooling Revolution

The transition from air cooling to liquid cooling is the most significant engineering shift in data centre design in a generation. Traditional air-cooled data centres — which use Computer Room Air Conditioning (CRAC) units and hot aisle/cold aisle containment to remove heat — are fundamentally limited in the power density they can support. At 10 kW per rack, air cooling is manageable; at 30 kW per rack, it becomes challenging; above 50 kW per rack, it is essentially impossible without liquid assistance.

Two primary liquid cooling technologies are scaling rapidly: Direct Liquid Cooling (DLC) — which pipes cooling liquid (typically water or a water-glycol mixture) directly to cold plates mounted on CPUs and GPUs — and immersion cooling, which submerges entire servers in thermally conductive dielectric fluid. Both approaches dramatically improve thermal management for high-density workloads, enable power densities of 100+ kW per rack, reduce PUE toward 1.0–1.05, and eliminate the need for large banks of air handlers and cooling towers (reducing water consumption).

The liquid cooling market is projected to grow at a CAGR exceeding 25% through 2030. By 2030, it is expected that the majority of new data centre construction will incorporate some form of liquid cooling, particularly for AI workloads. Retrofitting existing air-cooled facilities with liquid cooling infrastructure is technically complex and expensive (\$200–\$400 per kW), but increasingly necessary to serve the highest-value AI customers.

8.3 Edge Expansion & Distributed AI

The edge data centre market's growth trajectory — from \$7.2 billion in 2021 to a projected \$50–110 billion by 2030 at CAGRs of 16–25% — will be sustained and accelerated by the intersection of 5G network maturation and AI inference at the edge. As AI models become embedded in consumer applications, autonomous systems, industrial IoT, and smart city infrastructure, the need to run AI inference (applying trained models) closer to data sources will drive distributed edge AI deployments at unprecedented scale.

The combination of 5G's ultra-low latency and edge data centres' local compute creates a platform for applications that are impossible with centralized cloud alone: real-time autonomous vehicle decision-making, remote robotic surgery, AR/VR with imperceptible latency, smart manufacturing with real-time quality control at the production line level. These applications collectively represent a market of trillions of dollars in economic value — the data centre infrastructure serving them will grow commensurately.

The edge market is also being driven by data sovereignty concerns. Governments worldwide are enacting data localization laws requiring that citizen data be processed and stored domestically. This creates demand for edge data centres in markets that previously lacked meaningful digital infrastructure, from Southeast Asia to Africa to Latin America. The global total of edge deployments is expected to grow from a few hundred significant sites in 2023 to thousands by 2030.

8.4 Sustainability & Green Data Centres

The sustainability imperative will increasingly shape data centre investment decisions, design choices, and operational strategies through 2030 and beyond. Net-zero commitments from the world's largest hyperscalers — representing the majority of new data centre demand — are creating pull-through demand for renewable energy integration, green building certifications, water-efficient cooling, and carbon accounting transparency.

Wind and solar power purchase agreements (PPAs) have become standard financing tools for large data centre developments, with operators like Equinix and Microsoft routinely signing 200–500 MW renewable energy contracts to match their data centre consumption. Some hyperscalers are exploring more radical solutions: Amazon has announced investment in small modular nuclear reactors (SMRs) as a long-term source of carbon-free, always-on power for data centres. Microsoft has similarly explored nuclear power partnerships. The combination of renewable energy, SMRs, and high-efficiency

liquid cooling creates a pathway toward genuinely sustainable data centre operations, though the timeline and economics remain challenging.

The holistic sustainability measurement framework is also evolving: beyond PUE (Power Usage Effectiveness), the industry is adopting Water Usage Effectiveness (WUE) and Carbon Usage Effectiveness (CUE) as complementary metrics that together provide a more complete picture of data centre environmental performance. Regulatory reporting requirements — including the EU's Corporate Sustainability Reporting Directive (CSRD) and the SEC's proposed climate disclosure rules — are making sustainability data centre disclosures increasingly mandatory for large operators.

8.5 Market Projections to 2030 and 2035

Market Segment	2025E	2030P	2035P	CAGR
Global Data Centre Market	\$383B – \$430B	\$600B – \$700B	\$900B – \$1.1T	~11%
Edge Data Centre Market	\$14.7B – \$19.5B	\$50B – \$109B	\$140B – \$200B+	17–25%
AI Infrastructure (DC subset)	\$40B – \$60B	\$200B+	\$500B+	25%+
Data Centre Networking	\$45.8B	\$103B	\$200B+	17.6%
Liquid Cooling Market	\$~5B	\$20B+	\$50B+	25%+
U.S. Data Centre Market	\$134.8B	\$250B+	\$393B	11.3%
India Data Centre Market	\$11.5B	\$30B+	\$70B+	18%+
APAC Data Centre Capacity	12.2 GW	26+ GW	50+ GW	~15%

2030 Vision: The Data Centre Ecosystem

By 2030, the global data centre ecosystem will be characterized by: (1) Hyperscale AI factories of 500MW–1GW+ with full liquid cooling and on-site renewable/nuclear power; (2) A distributed edge network of thousands of sites providing sub-5ms latency to 90%+ of the world's population; (3) Fully digitized global supply chains connecting hyperscale cloud, enterprise hybrid cloud, and edge in a seamless multi-cloud fabric; (4) Sustainability-first design as a baseline regulatory requirement across all major markets; (5) A global data centre market approaching \$700B–\$1T in annual investment and activity.

9. CONCLUSIONS & STRATEGIC RECOMMENDATIONS

The global data centre industry has demonstrated exceptional resilience and growth across the 2020–2026 period, consistently outperforming broader economic growth and attracting record levels of institutional capital. The structural drivers — AI/GenAI, cloud computing, 5G/edge, and digital transformation — are each powerful independently, but their convergence creates a demand environment of extraordinary intensity.

For investors: The data centre sector offers compelling long-term returns backed by hard-asset security, contracted cash flows, and clear demand visibility. The highest risk-adjusted opportunities lie in pre-leased hyperscale development (certainty of demand) and in edge market pioneers (high growth potential, first-mover advantage). ESG-aligned facilities — those with renewable energy and liquid cooling from day one — will command premium rents and attract the most creditworthy tenants.

For operators: The critical success factors through 2030 are power access (securing grid capacity or on-site generation early), liquid cooling expertise (to serve AI workloads), sustainability credentials (to meet hyperscaler tenant requirements), and geographic diversification (to reduce dependence on any single constrained market). Operators without liquid cooling capabilities risk being excluded from the fastest-growing and highest-margin segment of the market.

For enterprises: The data centre landscape is bifurcating between public cloud (hyperscalers), private cloud (colocation and on-premise), and edge (distributed compute). The optimal strategy for most organizations is a hybrid multi-cloud architecture that leverages the strengths of each: scalable public cloud for dynamic workloads, secure private colocation for regulated or latency-sensitive applications, and edge compute for real-time, latency-critical use cases. Equinix's interconnection platform — connecting all three — is increasingly positioned as the neutral interconnection layer that makes this hybrid architecture work at global scale.

For policymakers: Data centre infrastructure has become critical national infrastructure. Governments that proactively facilitate data centre development — through clear permitting processes, grid investment, renewable energy support, and targeted incentives — will attract the investment that will underpin their digital economies for decades. Governments that restrict development through regulatory uncertainty, permitting delays, or unsustainable data sovereignty requirements will cede competitive advantage to more welcoming markets.

Industry Consensus: Key Metrics

Global Market 2025: ~\$383B – \$430B
 Global Market 2026E: ~\$425B – \$480B
 Market CAGR (2025–2033): ~11%
 Record Investment: \$61B+ (2025)

Forward-Looking Signals

Edge Market 2026E: \$16.9B – \$19.5B
 Edge CAGR 2026–2035: 17%–29%
 AI DC Investment to 2030: \$5T+
 Capacity by 2030: ~200 GW globally

Equinix 2026E Revenue: \$10.1B–\$10.2B
Equinix 2026E EBITDA Margin: ~51%
Vantage: \$25B+ Texas Mega-Campus

CapEx/MW (AI-ready): \$15M–\$20M+
Liquid Cooling CAGR: 25%+
Power Constraint: #1 bottleneck globally

— END OF REPORT —

Global Data Centre Industry Growth & Financial Analysis | April 2026

Sources: Grand View Research, Precedence Research, Equinix SEC Filings, Vantage Data Centres, Markets and Markets, Cushman & Wakefield, TrueLook, Thunder Said Energy, Nlyte, Alpha Matica

Researched & Curated by Charan GM | Civil Engineer | Data Centre Industry Research — Portfolio Project | April 2026